



Parallel DBSCAN on Spark

Group 12: Wang Jianming
Zhang Xinyue
Xing Zhenghao
Fu Chennan



2020/11/24

- **Introduction to DBSCAN**
- **Naive parallel DBSCAN**
 - Naive partition
 - Local DBSCAN
 - Merge
- **Optimization on partition**
 - Reduce-boundry based partition
 - Content based partition
- **Conclusion**





Local DBSCAN algorithm

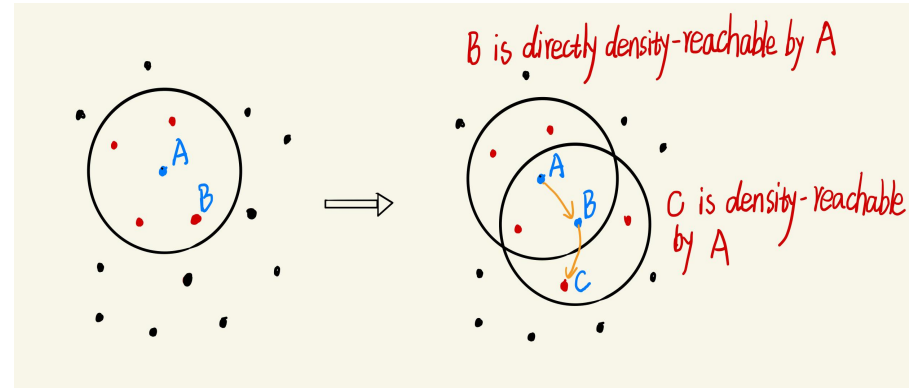
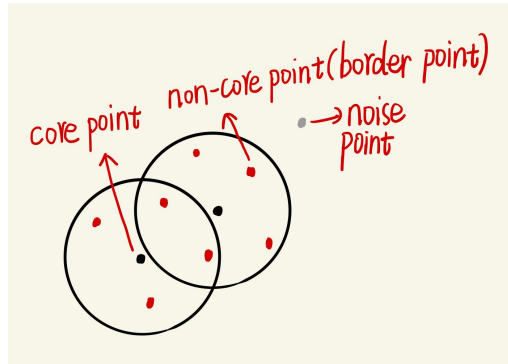


Concepts:

ϵ (Eps) — the given radius

MinPts — the given minimal points

if MinPts = 4:



Core idea:

1. select a data point p arbitrarily from the data set.
2. determine whether p is a core point by Eps and MinPts. If p is a core point, find all the data point which are **density-reachable from p** to form a cluster. If not, select another point to do the same steps as before.



Matrix DBSCAN algorithm

1. Introduction: Build a matrix to hold all those point-pair distance value instead of computing every time.
2. Pros: In principle, the computational cost could be reduced by half, which is not big algorithm-wise improvement, but a giant increase in engineering, especially it comes to big data problem.

example

	Point 1	2	3	4	5
point 1	X	$D_{1,2}$	$D_{1,3}$	$D_{1,4}$	$D_{1,5}$
2		X	$D_{2,3}$	$D_{2,4}$	$D_{2,5}$
3			X	$D_{3,4}$	$D_{3,5}$
4				X	$D_{4,5}$
5					X

→ Distance from P_1 to P_5



Parallel Matrix DBSCAN algorithm

Two Dimensional Example: (x, y)

Partition:

1. define partition number
2. define partition block position in the space
 - a. find the minimum x and y and find the maximum x and y
 - b. add eps to upper bound and minus eps to lower bound
 - c. divide the whole area into corresponding partition id (p_id)
3. build rdds (pid, dataset_points) by judging whether the points is in the partition area or n

Local DBSCAN:

1. each rdd will execute a matrix DBSCAN
2. each execution will return ((partition_dataset, core_point list), local_tags)

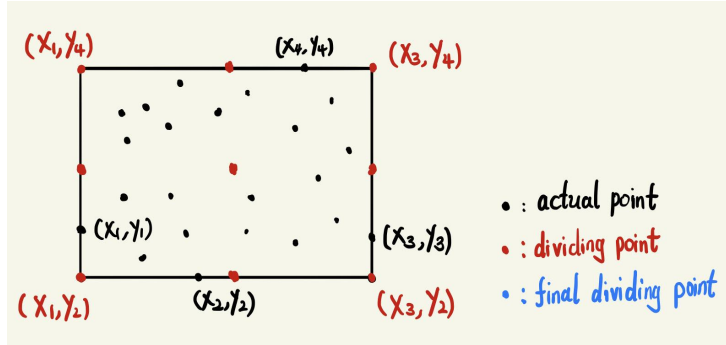
Merge:

1. update local tag list into global tag list
2. return final global tag list

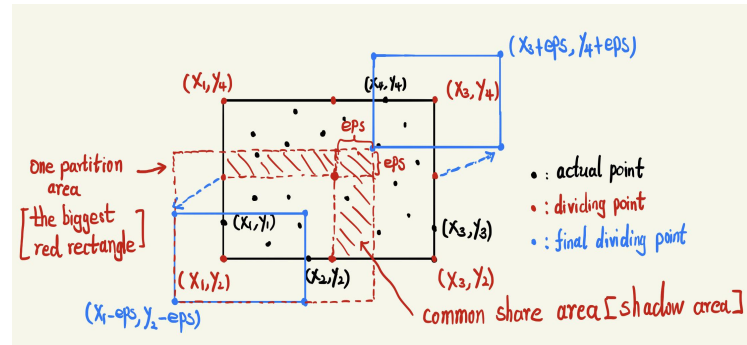
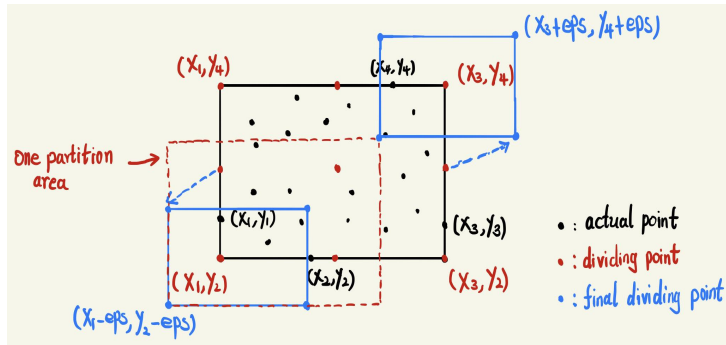


Parallel Matrix DBSCAN algorithm—Partition

partition number: 4 eps (epsilon): 3 MinPts (minimum points): 2



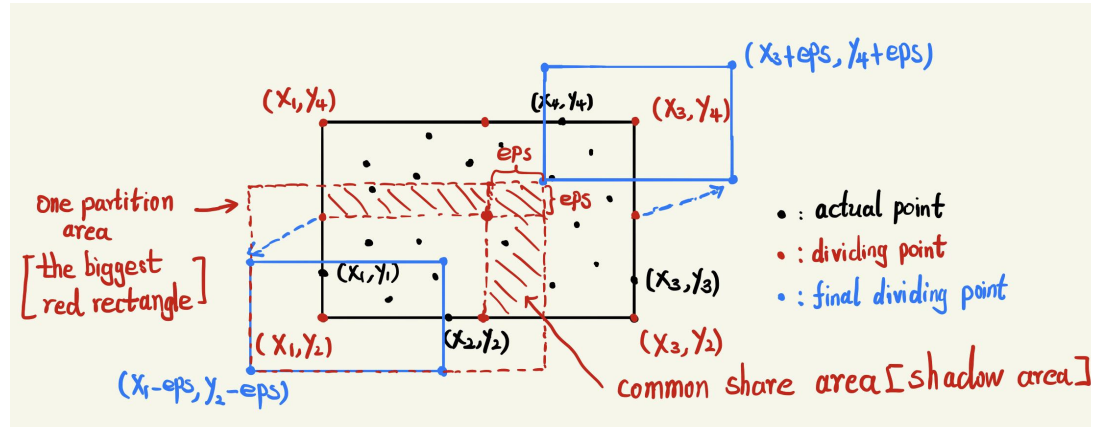
Define partition block position in the space



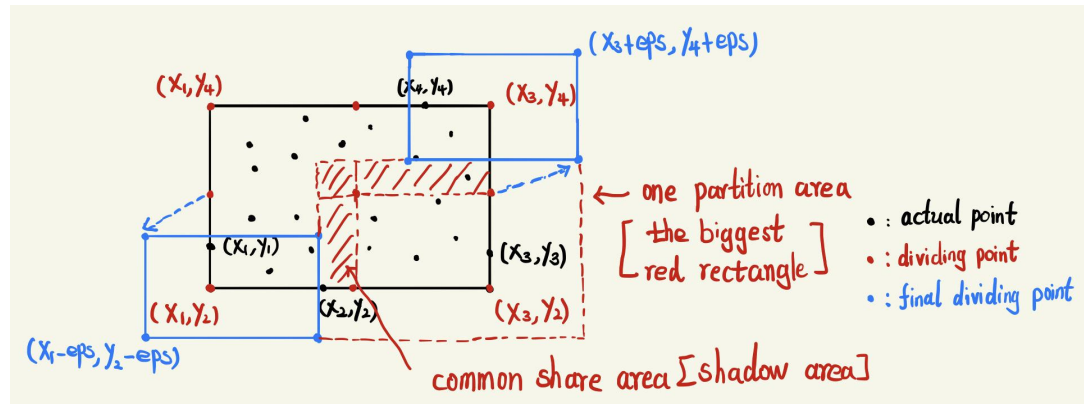


Parallel Matrix DBSCAN algorithm—Local Matrix DBSCAN

partition number: 4 eps (epsilon): 3 MinPts (minimum points): 2



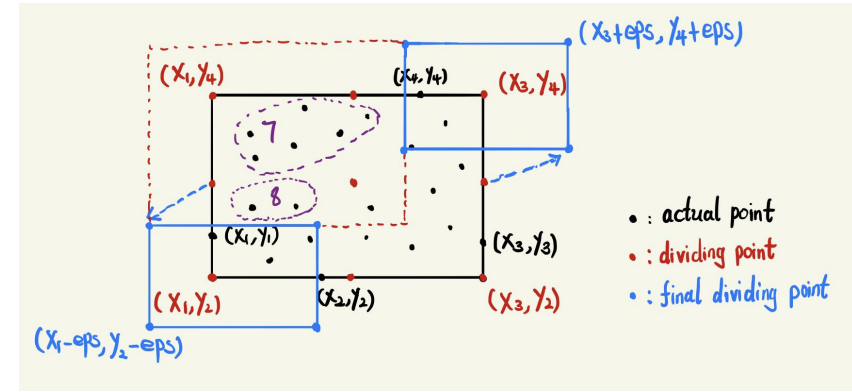
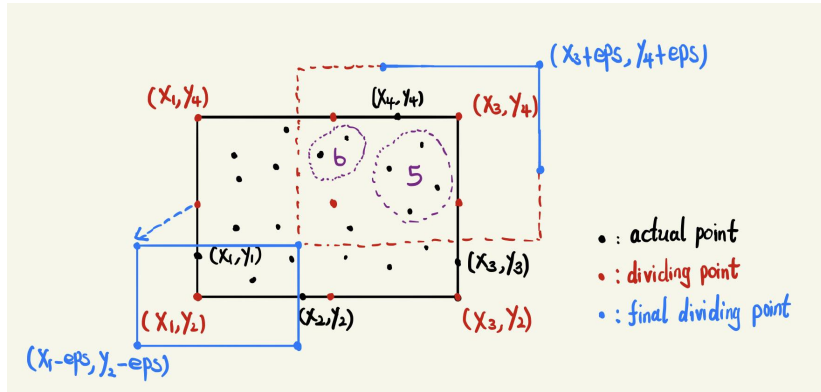
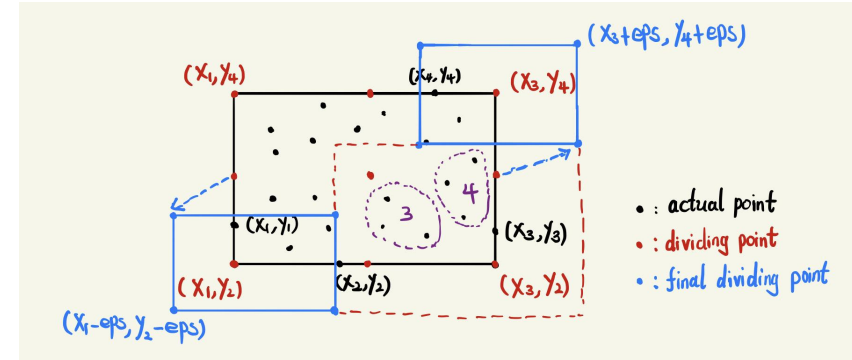
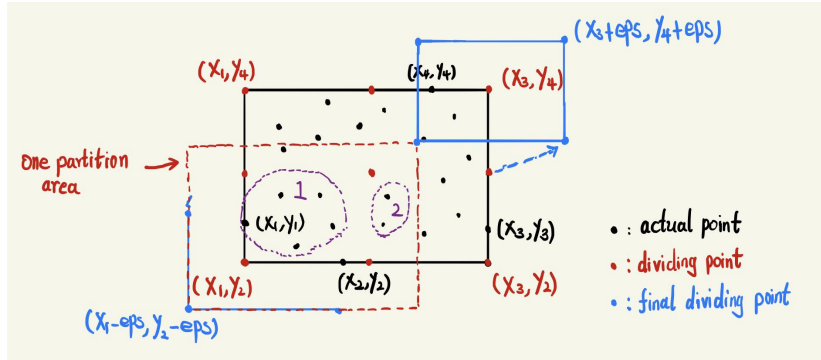
Points in the red shadow area will have multiple labels





Parallel Matrix DBSCAN algorithm—Local Matrix DBSCAN

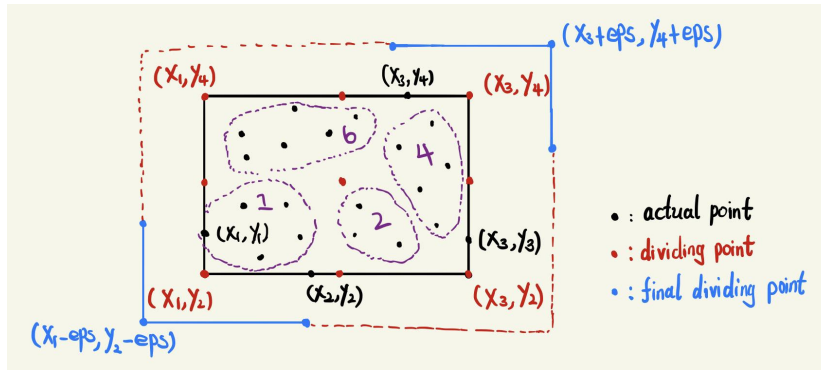
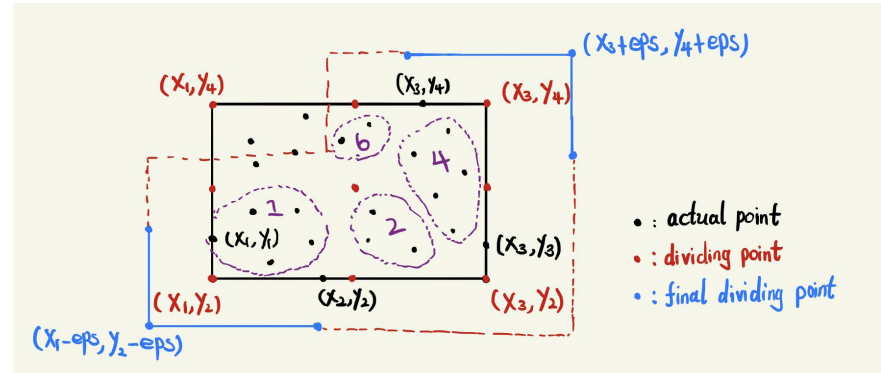
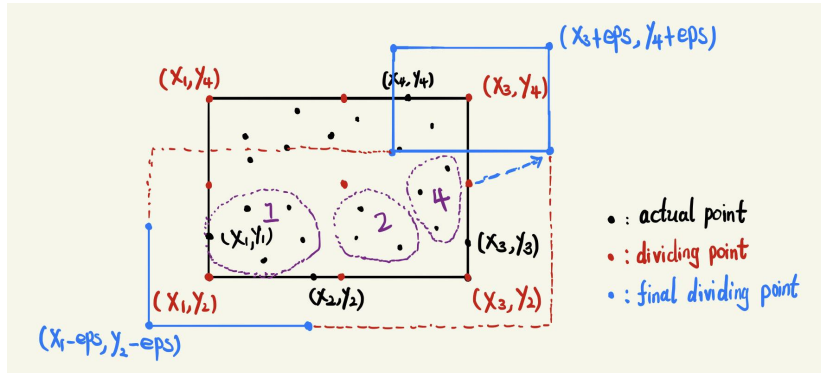
partition number: 4 eps (epsilon): 3 MinPts (minimum points): 2





Parallel Matrix DBSCAN algorithm—Merge

partition number: 4 eps (epsilon): 3 MinPts (minimum points): 2



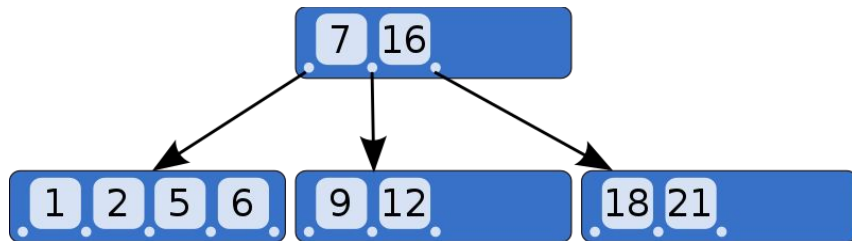
Using local labels to update global labels for all clusters

Optimization on partition——R tree



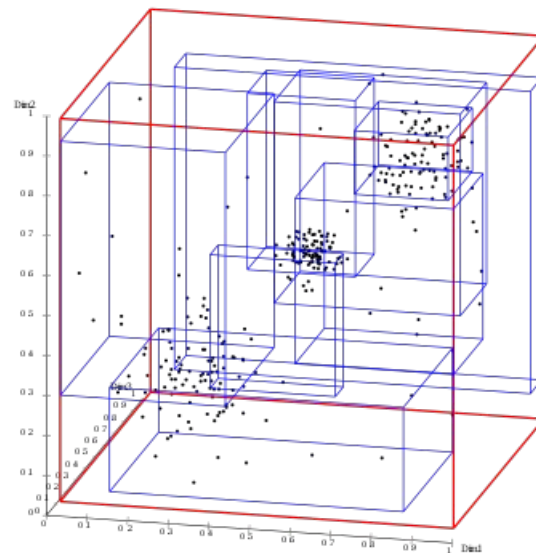
B-tree:

A self-balancing tree data structure

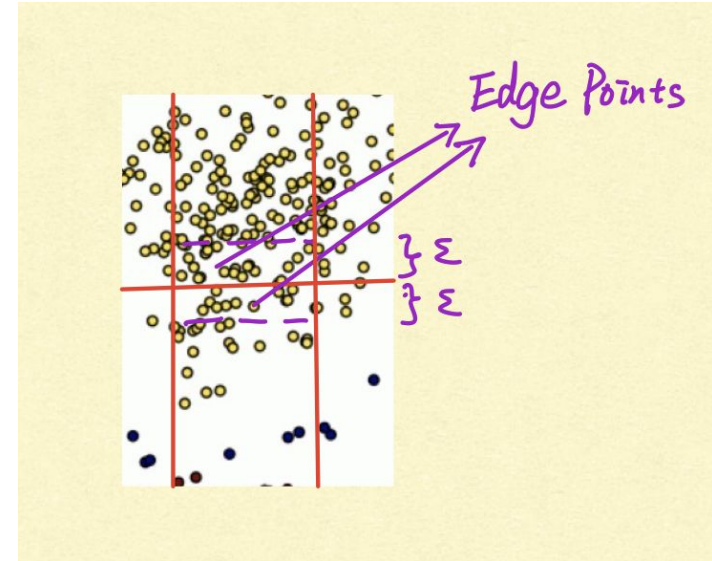
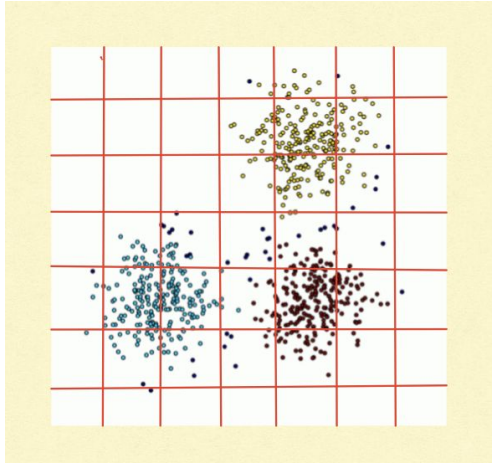


R-tree:

A tree data structures used for spatial access methods



Optimization on partition——R tree (Boundary-based Strategy)



Reduce the amount of edge points



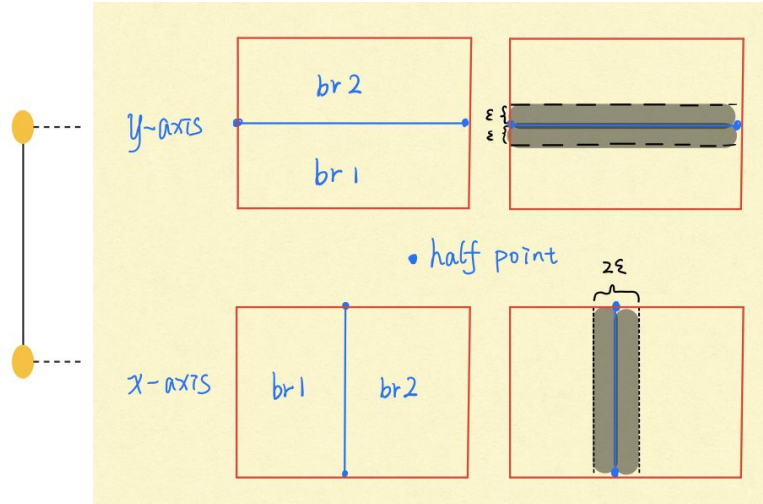
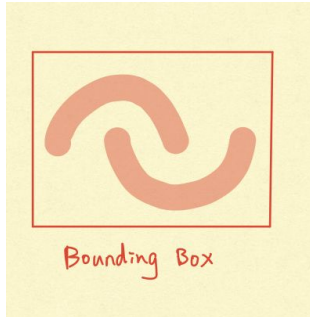
Make sure workload of each worker as balanced as possible



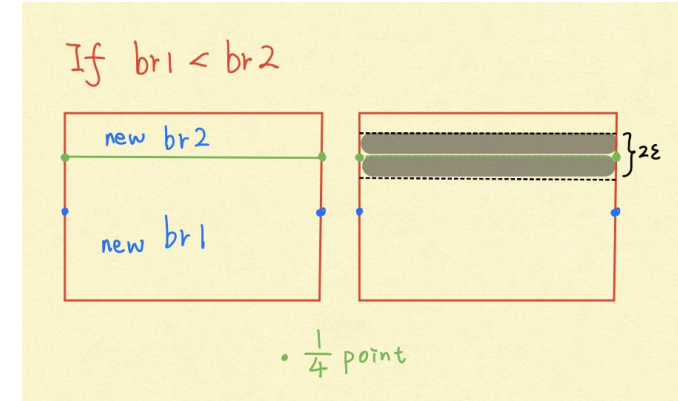
Optimization on partition—R tree (Boundary-based Strategy)



Set max points number in partition - maxpoints



As for y-axis:

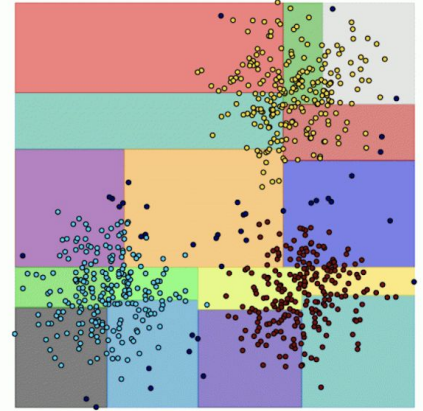
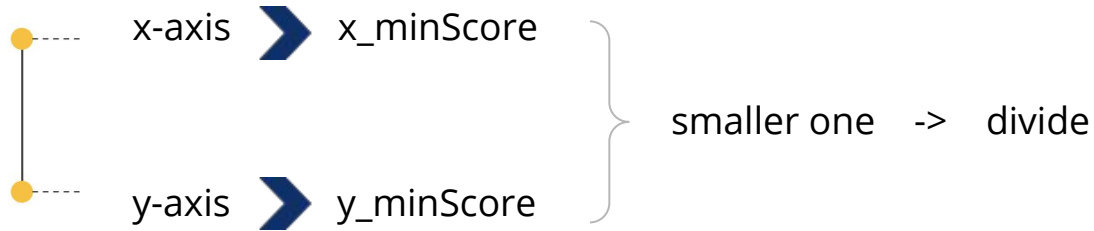


minimize **Score** = $|\# \text{ in br1} - \# \text{ in br2}| * (\# \text{ in Black part})$

Optimization on partition——R tree (Boundary-based Strategy)



Max points number in partition - maxpoints



Put final br1 & br2 to queue (right hand)

Take out the leftmost mbr in the queue in turn to do the above calculation



Optimization on partition—R tree (Cost-based Strategy)

Goal: Make the workload (# of points) of partitioned region as balanced as possible.

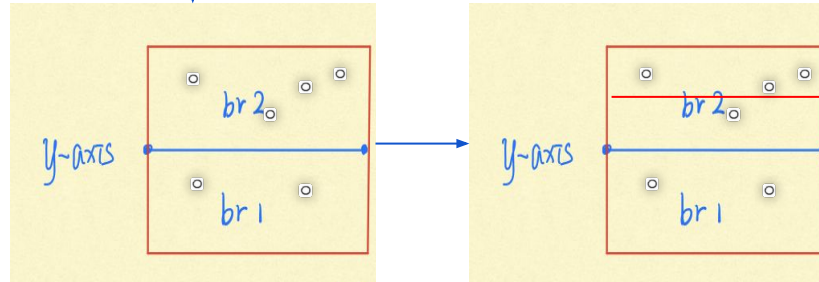
Input: **S** (The rectangle to be split)

Output: **S1** and **S2** such that **S1** \cup **S2** = **S**

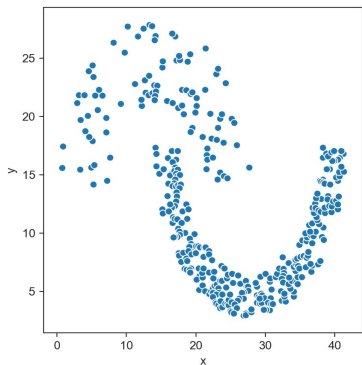
1. Min Cost Difference $\leftarrow \infty$

2. Calculate the cost difference between S1 and S2
(after applying the selected split line)

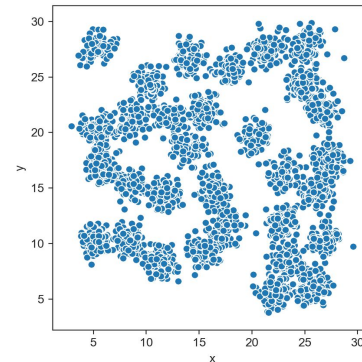
3. Split line + $\text{eps} * 2$ beyond the boundary
(Get the best split line from the record!)



Conclusion



Jain



D31

Partition Number	4	8	10	Partition Number	4	8	10
Naive	1179ms	797ms	808ms	Naive	26410ms	19415ms	18483ms
Parallel Matrix DBSCAN	422ms	385ms	380ms	Parallel Matrix DBSCAN	6730ms	5590ms	4564ms
Rtree - rbp	360ms	330ms	289ms	Rtree - rbp	7993ms	5722ms	5328ms
Rtree - cbp	308ms	284ms	271ms	Rtree - cbp	6977ms	5652ms	5164ms



Conclusion

- Performance
- Scalability
- Bottle neck (Master)
- Know your data

Future Work:

Now: using number of points to measure workload when doing partition

Future: using the real cost (maybe considering density) of DBSCAN to measure workload

More experiments on large volume datasets with large partitions

Partition Number	4	8	10	Partition Number	4	8	10
Naive	1179ms	797ms	808ms	Naive	26410ms	19415ms	18483ms
Parallel Matrix DBSCAN	422ms	385ms	380ms	Parallel Matrix DBSCAN	6730ms	5590ms	4564ms
Rtree - rbp	360ms	330ms	289ms	Rtree - rbp	7993ms	5722ms	5628ms
Rtree - cbp	308ms	284ms	271ms	Rtree - cbp	6977ms	5652ms	5264ms



Reference

- [1] DBSCAN Wikipedia. <https://en.wikipedia.org/wiki/DBSCAN>
- [2] Guttman, A. (1984). R-trees: a dynamic index structure for spatial searching. ACM SIGMOD International Conference on Management of Data.
- [3] Huang, Fang & Zhu, Qiang & Zhou, Ji & Tao, Jian & Zhou, Xiaocheng & Jin, Du & Tan, Xicheng & Wang, Lizhe. (2017). Research on the Parallelization of the DBSCAN Clustering Algorithm for Spatial Data Mining Based on the Spark Platform. Remote Sensing. 9. 10.3390/rs9121301.
- [4] He, Yaobin & Tan, Haoyu & Luo, Wuman & Feng, Shengzhong & Fan, Jianping. (2014). MR-DBSCAN: a scalable MapReduce-based DBSCAN algorithm for heavily skewed data. Frontiers of Computer Science. 8. 10.1007/s11704-013-3158-3.
- [5] Chernishev, G & Smirnov, K & Fedotovskiy, P & Erokhin, G & Cherednik, K. (2013). To Sort or not to Sort: The Evaluation of R-Tree and B+-Tree in Transactional Environment with Ordered Result Set Requirement. SYRCoDIS.
- [6] Clustering basic benchmark. <http://cs.joensuu.fi/sipu/datasets/>
- [7] Phnix-wu, Full analysis of spatial data index RTree, <https://blog.csdn.net/wzf1993/article/details/79547037>, 2018/03/12.